



Parallel Graph Algorithms in Intel Data Parallel C++ (DPC++)

Group 1

A bundle of glowing blue fiber optic cables, with many individual fibers visible, creating a starburst effect of light. The cables are arranged in a fan shape, pointing towards the right side of the frame. The background is dark blue and black.

Data Parallel C++ (DPC++)

DPC++


- Data Parallel C++
- New open source cross-architecture language
- Built upon ISO C++ and Khronos Group SYCL standards
- Part of a bigger oneAPI toolkit



Presenter: Rodger Chen

DPC++ Platforms

- DPC can be run on a variety of devices
 - CPU
 - GPU
 - FPGA
 - Accelerators



How does DPC work?

- An extension of SYCL
- SYCL is for portable heterogeneous C++ applications
- SYCL intended for openCL but is now more generic
 - Similarities between the two languages such as the same Host Kernel model.



Key Features:

- Single Source
- Host/Device model
- Asynchronous Execution


Single Source

- Single source file which contains both the host code and kernel code
- C++ based
 - No VHDL
- Easy to read and write - even for people that have basic coding knowledge

```
1  #include <CL/sycl.hpp>
2  using namespace sycl;
3
4  static const int N = 8;
5
6  int main(){
7      queue q;
8      std::cout << "Device: " << q.get_device().get_info<info::device::name>() << std::endl;
9
10     int *data = malloc_shared<int>(N, q);
11     for(int i=0; i<N; i++) data[i] = i;
12
13     q.parallel_for(range<1>(N), [=] (id<1> i){
14         data[i] *= 2;
15     }).wait();
16
17     for(int i=0; i<N; i++) std::cout << data[i] << std::endl;
18     free(data, q);
19     return 0;
20 }
21
```


Host / Kernel Model

- Host
 - Controls the execution of the program (“The brain”)
 - Usually the CPU, but does not have to be
- Kernel
 - Refers to the code that will be run on the device
 - Can be run on a number of platforms



Asynchronous Execution

- Any work sent to a device is done asynchronously from the Host.
- Amdahl's law!
- If implemented correctly: scalable
- If implemented incorrectly: Race conditions
- Data dependencies can be satisfied in a multitude of ways in DPC++


```
// Task A
auto eA = Q.submit([&](handler &h) {
    h.parallel_for(N, [=](id<1> i) { /*...*/ });
});
eA.wait();

// Task B
auto eB = Q.submit([&](handler &h) {
    h.parallel_for(N, [=](id<1> i) { /*...*/ });
});

// Task C
auto eC = Q.submit([&](handler &h) {
    h.depends_on(eB);
    h.parallel_for(N, [=](id<1> i) { /*...*/ });
});


// Task D
auto eD = Q.submit([&](handler &h) {
    h.depends_on({eB, eC});
    h.parallel_for(N, [=](id<1> i) { /*...*/ });
});

return 0;
```



DPC++ Language Features

- Similar to OpenCL
- Queues
- Buffers
- Kernels




```
// Create buffers that hold the data shared between the host and the devices.  
// The buffer destructor is responsible to copy the data back to host when it  
// goes out of scope.  
buffer a_buf(a_vector);  
buffer b_buf(b_vector);  
buffer sum_buf(sum_parallel.data(), num_items);
```

```
// Submit a command group to the queue by a lambda function that contains the  
// data access permission and device computation (kernel).
```

```
q.submit([&](handler &h) {  
    // Create an accessor for each buffer with access permission: read, write or  
    // read/write. The accessor is a mean to access the memory in the buffer.  
    accessor a(a_buf, h, read_only);  
    accessor b(b_buf, h, read_only);  
  
    // The sum_accessor is used to store (with write permission) the sum data.  
    accessor sum(sum_buf, h, write_only, noint);  
  
    // Use parallel_for to run vector addition in parallel on device. This  
    // executes the kernel.  
    // 1st parameter is the number of work items.  
    // 2nd parameter is the kernel, a lambda that specifies what to do per  
    // work item. The parameter of the lambda is the work item id.  
    // DPC++ supports unnamed lambda kernel by default.  
  
    h.parallel_for(num_items, [=](auto i) { sum[i] = a[i] + b[i]; });  
});
```

Queues

- Abstraction for connecting to a device
- Bound to a single device at a time
- Queue has a submit method which takes a command group
- Work is transferred to the Device




```
queue q{}; //Default to most performant device
queue q_CPU(host_selector{}, exception_handler);
queue q_GPU(gpu_selector{}, exception_handler);
```

Queues

- Abstraction for connecting to a device
- Bound to a single device at a time
- Queue has a submit method which takes a command group
- Work is transferred to the Device

```
q.submit([&](handler &h) {  
    // Create an accessor for each buffer with access permission: read, write or  
    // read/write. The accessor is a mean to access the memory in the buffer.  
    accessor a(a_buf, h, read_only);  
    accessor b(b_buf, h, read_only);  
  
    // The sum_accessor is used to store (with write permission) the sum data.  
    accessor sum(sum_buf, h, write_only, noinit);  
  
    h.parallel_for(num_items, [=](auto i) { sum[i] = a[i] + b[i]; });  
}
```

```
// Task A  
auto eA = Q.submit([&](handler &h) {  
    h.parallel_for(N, [=](id<1> i) { /*...*/ });  
});  
eA.wait();  
  
// Task B  
auto eB = Q.submit([&](handler &h) {  
    h.parallel_for(N, [=](id<1> i) { /*...*/ });  
});  
  
// Task C  
auto eC = Q.submit([&](handler &h) {  
    h.depends_on(eB);  
    h.parallel_for(N, [=](id<1> i) { /*...*/ });  
});  
  
// Task D  
auto eD = Q.submit([&](handler &h) {  
    h.depends_on({eB, eC});  
    h.parallel_for(N, [=](id<1> i) { /*...*/ });  
});  
  
return 0;
```



```
#include <CL/sycl.hpp>
#include <array>
#include <iostream>
using namespace sycl;

int main() {
    constexpr int size=16;
    std::array<int, size> data;

    // Create queue on implementation-chosen default device
    queue Q;

    // Create buffer using host allocated "data" array
    buffer B { data };

    Q.submit([&](handler& h) {
        accessor A{B, h};
        h.parallel_for(size, [=](auto& idx) {
            A[idx] = idx;
        });
    });

    // Obtain access to buffer on the host
    // Will wait for device kernel to execute to generate data
    host_accessor A{B};
    for (int i = 0; i < size; i++)
        std::cout << "data[" << i << "] = " << A[i] << "\n";


    return 0;
}
```

The diagram illustrates the organization of the code into three sections:

- Host code**: Indicated by a blue bracket on the right, it encompasses the code from the beginning of the `main` function up to the `Q.submit` call.
- Device code**: Indicated by a red bracket on the right, it encompasses the code within the `Q.submit` block, which is also enclosed by a red dashed line.
- Host code**: Indicated by a blue bracket on the right, it encompasses the code following the `Q.submit` block, including the `host_accessor` and the `for` loop.

Memory Management

- Buffers and Accessors
 - Buffers for holding data
 - Memory can be located anywhere
 - Accessors for accessing
- Unified Shared Memory
 - Familiar C++ Pointer Style
 - Implicit or explicit data transfer



Allocation Type	Description	Accessible on host?	Accessible on device?	Located on
device	Allocations in device memory	✗	✓	device
host	Allocations in host memory	✓	✓	host
shared	Allocations shared between host and device	✓	✓	can migrate back and forth

Memory Management

- Buffers and Accessors
 - Buffers for holding data
 - Memory can be located anywhere
 - Accessors for accessing
- Unified Shared Memory
 - Familiar C++ Pointer Style
 - Implicit or explicit data transfer

```
// Create buffers that hold the data shared between the host and the devices.  
// The buffer destructor is responsible to copy the data back to host when it  
// goes out of scope.  
buffer a_buf(a_vector);  
buffer b_buf(b_vector);  
buffer sum_buf(sum_parallel.data(), num_items);
```

```
// Create an accessor for each buffer with access permission: read, write or  
// read/write. The accessor is a mean to access the memory in the buffer.  
accessor a(a_buf, h, read_only);  
accessor b(b_buf, h, read_only);  
  
// The sum_accessor is used to store (with write permission) the sum data.  
accessor sum(sum_buf, h, write_only, noinit);
```

Explicit memory operation	copy	Copy data between locations specified by accessor, pointer, and/or shared_ptr. The copy occurs as part of the DAG, including dependence tracking.
	update_host	Trigger update of host data backing of a buffer object.
	fill	Initialize data in a buffer to a specified value.

Kernels

- Code which runs on the device
- Platform Agnostic
 - Describes parallelism
- Lambda Function
 - Capture
 - Parameters
 - Lambda Body

```
// Use parallel_for to run vector addition in parallel on device. This
// executes the kernel.
// 1st parameter is the number of work items.
// 2nd parameter is the kernel, a lambda that specifies what to do per
// work item. The parameter of the lambda is the work item id.
// DPC++ supports unnamed lambda kernel by default.

h.parallel_for(num_items, [=](auto i) { sum[i] = a[i] + b[i]; });
```

```
[ capture-list ] ( params ) -> ret { body }
```

Actions Types

- Single task
- Parallel for
 - Basic Kernel
- Parallel for work group
 - Thread of execution
 - ND-Kernel
 - Hierarchical kernel

Work Type	Actions (handler class methods)	Summary
Device code execution	single_task	Execute a single instance of a device function.
	parallel_for	Multiple forms are available to launch device code with different combinations of work sizes.
	parallel_for_work_group	Launch a kernel using hierarchical parallelism, described in Chapter 4.
Explicit memory operation	copy	Copy data between locations specified by accessor, pointer, and/or shared_ptr. The copy occurs as part of the DAG, including dependence tracking.
	update_host	Trigger update of host data backing of a buffer object.
	fill	Initialize data in a buffer to a specified value.


Parallel_for	#pragma omp parallel for
Single_task ~~	#pragma omp critical

Graph Algorithms




What is a Graph?

- Composed of vertices/nodes and edges
- Edges are used to represent relation between vertices
- Weighted vs Unweighted
- Directed vs Undirected
- Cyclic vs Acyclic



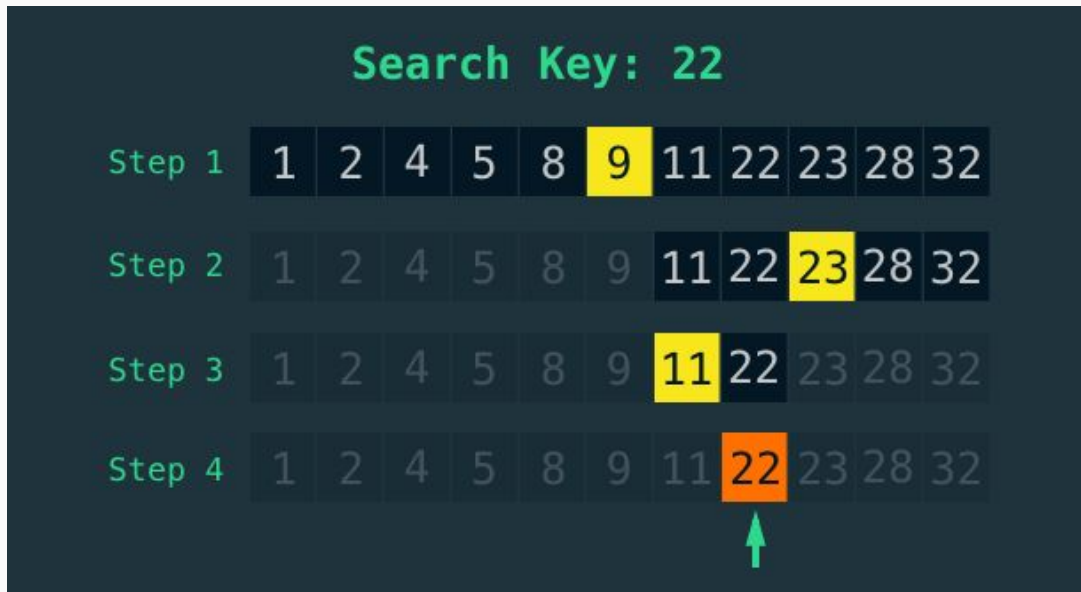
Tree



Forest

Binary Search

- Used to find (search for) a specific node
- Used on sorted arrays
- Continues to compare middle value with target value until target is found
- Useful for going through sorted databases



Binary Search

- Naive implementation: Split nodes based on number of processors and perform binary search on each of them
- Parallelize the comparisons by splitting into p groups instead of 2
- Parallelizing multiple binary searches on the same tree


```
parallel for (number of subtrees){
    PerformBinarySearch(subtree, endnode);
}

parallel for (nodes connecting to current node){
    if (the new node is the target node){
        return target node has been found;
    }
    else if (distance from new node to target node is less than the previous closest node) {
        closest node = new node;
    }
    return closest node;
}

parallel for (number of nodes that need to be found){
    PerformBinarySearch(tree, endnode);
}
```


Breadth First Search

- Graph traversal algorithm
- Simple
- Used as the foundation for more complicated algorithms
- Useful in many real-world applications



Breadth First Search


- Process each depth level in parallel
- Each processor assigned a node, searches for children
- Critical sections are needed when queuing and dequeuing
- Pseudocode in OpenMP, we will translate to DPC++

```
while (!q.empty()) {
    qSize = q.size();

    // process entire level in parallel
    #pragma omp parallel for
    for (int i = 0; i < qSize; i++) {
        node* currNode;
        #pragma omp critical
        {
            currNode = q.front();
            q.pop();
        }
        // find children for the current node
        node* childNodes;
        findChildren(currNode);
        #pragma omp critical
        q.push(childNodes);
    }
}
```

Prim's Algorithm

- An algorithm that finds the minimum spanning tree on a graph
- Greedy algorithm, always makes the local optimal choice
- Simple
- Useful for telecommunications or solving complex mathematical problems



Prim's Algorithm

- Main loop not parallelizable; however, inner loop is
- Slow for large numbers of nodes, we can use parallelisation to overcome this issue
- Current parallel implementation in OpenMP, we will translate to DPC++

```
// process inner loop in parallel
#pragma omp parallel for
for(j = 0; j < prim.dim; j++) {
    // find the minimum weight
    if(prim.edge[prim.U[i]][j] > minDist || prim.edge[prim.U[i]][j]==0) {
        continue;
    } else {
        #pragma omp critical
        {
            h.parallel_for(nd_range{{size}, {16}}, [=](nd_item<1> item) {
                auto sg = item.get_sub_group();
                ...
                sg.barrier();
                ...
            });
        }
    }
}
```

Questions?